



Trinity College Dublin

Coláiste na Tríonóide, Baile Átha Cliath

The University of Dublin

SCHOOL OF COMPUTER SCIENCE AND STATISTICS

MODERN DELAY-TOLERANT EMAIL

RAKESH LAKSHMANAN

STEPHEN FARRELL

AUGUST 5, 2026

SUBMITTED IN PARTIAL FULFILMENT OF THE REQUIREMENTS FOR THE DEGREE OF
M.SC. COMPUTER SCIENCE - FUTURE NETWORKED SYSTEMS

Abstract

Communications links in deep space suffer from high delays and disruptions that make protocols designed for a low-latency, always-available network unsuitable for use over them. Delay-Tolerant Networking (DTN) allows e-mail to be delivered over such links, but the problem of filtering spam in that environment has been little studied. This matters because many widely-used spam filters depend on live lookups to remote servers, which must now be reached over the delayed link. This research measures the performance of three widely-used filters, SpamAssassin, rspamd, and Vipul's Razor, when they are accessed under high delay, and draws conclusions and recommendations from the results.

A testbed with five nodes was created to test an emulated Earth–Moon link, consisting of real applications: Postfix and Dovecot for message delivery, ION-DTN and BMail gateway for message transmission, and a locally installed DNS resolver for blocklist lookups. A delay equal to 1.3 seconds was imposed on the relay node and each filter was tested using 681 known-spam messages under three scenarios, a no-delay baseline, the delay applied with the tool's default timeouts, and the delay with the tool's own timeout reduced below the round-trip, while recording per-message verdicts, scan times, and packet captures.

Acknowledgements

I would like to thank my supervisor, Professor Stephen Farrell. Under his supervision, I was able to successfully complete this dissertation. His inspiration and suggestions regarding the project's direction were essential for the progress of this research. I am also grateful to Richard Clayton for providing the spam message corpus used in this evaluation.

Finally, I would like to thank my family and friends for their support and encouragement throughout this process.

Contents

Abstract	i
1 Introduction	1
1.1 Background	1
1.2 Motivation	2
1.3 Research Question and Objectives	3
1.4 Dissertation Overview	4
2 Literature Review	5
2.1 Background	5
2.1.1 Delay-Tolerant Networking and the Bundle Protocol	5
2.1.2 ION	6
2.2 Email over DTN	6
2.3 Anti-spam filtering	8
2.4 DTN security	8
2.5 Related problems	9
2.6 Methods and their limits	10
2.7 The gap	11
2.8 Summary	11
3 System Design	12
3.1 System Architecture	12
3.2 Environment and Technology Selection	14

3.2.1	Virtual machines over containers	14
3.3	Email stack	15
3.4	DTN stack	15
3.5	Anti-Spam Tool Selection	15
3.6	DTN Design	16
3.7	Delay Simulation Design	16
3.7.1	DNS Resolver	17
3.8	Summary	18
4	Implementation	19
4.1	Basis and Scope of This Work	19
4.2	Environment and Dependencies	19
4.3	DTN Layer	20
4.4	Mail Layer	21
4.4.1	Build	22
4.4.2	DTPC configuration	22
4.4.3	Postfix and Dovecot	23
4.5	Delay Simulation	23
4.6	DNS Resolver	24
4.7	Anti-Spam Tools	25
4.7.1	SpamAssassin	25
4.7.2	rspamd	26
4.7.3	Vipul's Razor	26
4.8	Email Corpus	27
4.9	Run Harness	27
4.9.1	Structure	28
4.9.2	Preflight and boot	28
4.9.3	Checks before a run	29
4.9.4	Sequence of a run	29
4.9.5	Completion and collection	31

4.10	Running an Experiment	32
4.11	Summary	32
5	Results and Discussion	34
5.1	SpamAssassin	34
5.1.1	Detection	34
5.1.2	Scan time	36
5.1.3	Flips	37
5.1.4	Network volume	38
5.2	rspamd	38
5.2.1	Detection	39
5.2.2	Scan time	39
5.2.3	Flips, Run 2 to Run 3	40
5.2.4	Network volume	40
5.3	Vipul's Razor	41
5.3.1	Detection	41
5.3.2	Scan time	42
5.3.3	Flips	42
5.3.4	Network volume	43
5.4	Summary	44
6	Conclusions and Future Work	47
6.1	Conclusion	47
6.2	Limitations	48
6.3	Future Work	49
A1	Appendix	56
A1.1	Use of Artificial Intelligence	56

List of Figures

3.1	Testbed topology. Each client connects only to its local mail server. The relay is the only path between the Earth and Moon networks and is where the one-way link delay is applied.	13
3.2	End-to-end path of an email from Earth to Moon across the DTN link. .	14
3.3	Symmetric delay. Each leg crossing the relay is delayed by 1300 ms: the Earth-to-Moon leg on the interface's egress path, and the Moon-to-Earth leg on an ifb device, since ingress cannot be shaped directly. The two legs together give a round-trip time of about 2.6 s. Traffic that stays within one network does not cross the relay and is not delayed.	17

List of Tables

5.1	SpamAssassin - Aggregate results across the three runs.	34
5.2	SpamAssassin - DNSBL/URIBL rule firing (selected rules); every rule falls from Run 2 to Run 3.	35
5.3	SpamAssassin - DNSBL-hit rate by corpus position; Run 3 collapses after $\approx$ email 500.	36
5.4	SpamAssassin - Scan-time distribution (emails per bucket).	37
5.5	SpamAssassin - Verdict changes and blocklist drift, Run 2 to Run 3. . . .	37
5.6	SpamAssassin - Traffic on the relay: DNS to the resolver vs DTN bundle traffic.	38
5.7	rspamd - Detection and scoring across the three runs.	39
5.8	rspamd - Scan-time distribution across the three runs (messages per bucket).	39
5.9	rspamd - Per-message flips between the two delayed runs.	40
5.10	rspamd - Network volume per run.	40
5.11	Razor - Verdict distribution across the three runs.	41
5.12	Razor - Scan-time distribution across the three runs.	42
5.13	Razor - Per-message flips, matched on a Message-ID and Delivery-date key.	42
5.14	Razor - Network volume per run.	43
5.15	Summary of tool behaviour under delay.	44
5.16	Timeout settings that govern behaviour under delay, and the effect of setting them relative to the round-trip time ($RTT \approx 2.6$ s).	45

5.17	Summary of network traffic volume across all tools. SpamAssassin reports DNS traffic only, rspamd reports total captured traffic, and Razor reports total captured traffic at the receiver.	46
5.18	Behaviour common to SpamAssassin, rspamd, and Razor under delay. .	46

List of Abbreviations

Abbreviation	Expansion
ARM64	64-bit ARM Architecture
BP	Bundle Protocol
DKIM	DomainKeys Identified Mail
DMARC	Domain-based Message Authentication, Reporting, and Conformance
DNS	Domain Name System
DNSBL	DNS-based Blocklist
DNSSEC	Domain Name System Security Extensions
DTN	Delay-Tolerant Networking
DTPC	Delay-Tolerant Payload Conditioning
HTTP	Hypertext Transfer Protocol
IETF	Internet Engineering Task Force
ifb	Intermediate Functional Block
IMAP	Internet Message Access Protocol
ION	Interplanetary Overlay Network
LTS	Long-Term Support
RTO	Retransmission Timeout
RTT	Round-Trip Time
SDR	Simple Data Recorder (ION persistent store)
SMTP	Simple Mail Transfer Protocol
SPF	Sender Policy Framework
SSH	Secure Shell
SYN	Synchronize (TCP segment)
TCP	Transmission Control Protocol
TLS	Transport Layer Security
UDP	User Datagram Protocol
URI	Uniform Resource Identifier
URIBL	URI-based Blocklist
VM	Virtual Machine

1 | Introduction

1.1 Background

Communication in space is not like communication on Earth. The distances are large, so signals take time to travel: a one-way trip between Earth and the Moon takes about 1.3 seconds, and to Mars it takes minutes. Links are also interrupted for long periods, as bodies move and block the line of sight. The normal internet assumes the opposite, short delays and a connection that stays up, so its protocols may fail here: they expect a reply within a short time and treat a missing one as an error.

Delay-Tolerant Networking (DTN) (1) was designed for these conditions. Instead of a live end-to-end connection, it uses store-and-forward: each node holds a message until the next link is available, then passes it on. The Bundle Protocol (2) carries messages as self-contained units called bundles between DTN nodes. This tolerates long delays and interrupted links, which makes DTN the accepted basis for interplanetary networking (3), an active area with real implementations (4, 5). Space agencies are already building detailed plans for lunar exploration on DTN and related protocols (6)

Email is one of the basic tools people need, so carrying it over DTN has received attention. In one approach, a gateway translates between standard mail protocols and the Bundle Protocol (7, 8), and open tools implement this (9). The mail servers act as the mail transfer agents (MTAs) on each side, and the Bundle Protocol carries messages between them, so a message can move from a server on one side of a DTN

link to the other.

On any mail system, much of the incoming mail is spam, and anti-spam tools are a standard part of the setup. These tools do more than read the message text. They also analyse the sending MTA and the links in the message against live network services: DNS blocklists of known-bad addresses, sender reputation systems, and shared databases of known spam. Many of these checks send a request to a remote server during the scan and wait for an answer before they can score that rule.

1.2 Motivation

Work on DTN email has gone into moving mail across the link (7, 8, 9), not into filtering it once it arrives. Yet filtering is part of running a mail service, not an optional extra. Consider a crewed Moon base whose users exchange email with Earth over DTN (6): spam will reach it as it reaches any mail system, and the tools that would catch it now depend on lookups that must cross the delayed link. Filtering probably cannot simply be left to Earth. More than one space agency may be involved, and mail for one agency's crew can arrive through another's infrastructure, so an operator may not trust Earth-side filtering it does not control and will want to run its own anti-spam at the point of delivery, on the Moon side.

Whether those tools still work is not obvious. DTN breaks an assumption they rely on: an answer is now seconds away, or the link is down when a request is made. A tool might scan far more slowly while it waits, miss spam when answers arrive too late, give up early and let spam through, retry until it overloads the link, or fail with no sign that it failed. Each tool reaches the network in its own way, so each may respond differently. To our knowledge, none of this has been measured in a scenario like this. The value is in understanding how the tools behave under delay and why, which is a starting point for anyone who later builds spam filtering into interplanetary email.

1.3 Research Question and Objectives

This dissertation turns that gap into a precise question.

Research question. The primary research question is:

How do widely-used anti-spam tools behave when run over a delay-tolerant network link, and what mechanisms drive that behaviour?

It breaks into three sub-questions:

- How does each tool's network dependence interact with propagation delay?
- What failure modes appear under delay, and what causes them?
- How do the tools' own timeout and retry settings behave under delay, and does tuning them change the outcome?

Objectives. To answer this, the dissertation sets out to:

1. Build a controlled Earth–Moon DTN email testbed using widely-used mail and DTN software, with a link delay that can be set to a chosen value.
2. Use a corpus of recent spam messages and a repeatable method to scan it and record per-email timing, verdicts, and network traffic.
3. Evaluate three tools, SpamAssassin, rspamd and Vipul's Razor, chosen to cover the main ways a filter depends on the network.
4. For each tool, compare its behaviour under delay against a no-delay baseline and identify the mechanisms behind the differences and any failure modes.
5. Draw the cross-tool comparison: what generalises about anti-spam filtering that depends on synchronous network services when it runs under delay.

Scope. This is a behavioural study: it characterises and explains the tools' response to delay, and does not propose or build a filtering system for interplanetary

email.

Contributions. The dissertation makes two contributions:

- A reproducible Earth–Moon DTN email testbed for evaluating mail-layer tools under controlled delay.
- An empirical, mechanism-level account of how three anti-spam tools behave under delay, including the failure modes and their causes.

1.4 Dissertation Overview

This dissertation is organised into six chapters. Chapter 1 introduces the problem, the research question, and the objectives. Chapter 2 reviews the background:

delay-tolerant networking and the Bundle Protocol, the email system and how mail is carried over DTN, and the anti-spam tools and the network services they depend on. Chapter 3 describes the design of the testbed, including the node layout, the choice of software, and the delay simulation. Chapter 4 covers the implementation, from building the nodes to configuring the mail, DTN, and anti-spam software.

Chapter 5 presents the evaluation: each of the three tools is measured under delay against a no-delay baseline, and the results are compared. Chapter 6 concludes and sets out directions for future work.

2 | Literature Review

This chapter reviews the work that surrounds spam filtering over DTN. It covers three fields that touch the subject, a set of problems elsewhere that have the same shape, and the methods those fields use to evaluate their systems. It ends by identifying the gap this dissertation addresses: the behaviour of anti-spam filtering under delay.

2.1 Background

This section introduces the parts of delay-tolerant networking that the later chapters rely on: the Bundle Protocol and its terms, and ION, the implementation used to build the testbed.

2.1.1 Delay-Tolerant Networking and the Bundle Protocol

A node is a single participant in a delay-tolerant network. A bundle is the self-contained unit the Bundle Protocol carries between nodes (2), holding the message together with the information needed to forward it. Under the ipn naming scheme, each node has a number, and an endpoint on that node is written `ipn:node.service`, where the service number identifies the application that receives the bundle.

The Bundle Protocol does not itself move bundles between neighbouring nodes. That is left to a convergence layer, an adapter that carries bundles over an underlying

transport such as UDP or TCP. The testbed uses the UDP convergence layer, so each bundle travels between nodes inside a UDP datagram.

2.1.2 ION

ION, the Interplanetary Overlay Network, is one implementation of the Bundle Protocol, developed at NASA's Jet Propulsion Laboratory (4). It was chosen for its use in deep-space missions and for the mail tools built on it. The paragraphs below cover the parts of ION the later chapters depend on: how it stores bundles, how it routes them, and the delivery service the mail gateway uses.

Because a delay-tolerant network holds a message until the next link is available, each node needs somewhere to keep it in the meantime. In ION this is the SDR, a persistent store where bundles and protocol state are held so they survive a break in the link or a restart of the node.

A deep-space link is also not always available, so ION does not discover routes on demand. It computes them from a contact plan, a schedule that declares when the link between two nodes is up and at what data rate, together with the propagation delay between them. From this plan, ION's Contact Graph Routing works out how a bundle should reach its destination.

Above the Bundle Protocol, ION offers Delay-Tolerant Payload Conditioning (DTPC), a service that provides reliable, in-order delivery and can gather several application messages into one bundle. The mail gateway sends over DTPC rather than over raw bundles.

2.2 Email over DTN

Transporting email over a delayed or intermittent network connection is not novel. Previous studies are terrestrial and were developed for rural regions with no continuous connection. DakNet (10) used automobiles to transport email from

village to village, and to the internet gateway, offering an inexpensive alternative to a wired connection. KioskNet (11) improved on this with a whole system including kiosk controllers, naming, routing and a security framework, reporting a pilot installation. MotoPost (12) used dual-mode mobile phones as data mules, with SMTP as the base application protocol for message transfer.

In all three systems, security means authenticating users and keeping message contents confidential, rather than filtering the mail itself. KioskNet (11) uses a public-key-infrastructure framework to provide privacy, confidentiality and integrity, and handles user management and logging. None of the three has filters to remove junk mail, and none reports how much was received during their deployments.

A more recent set of IETF drafts addresses email over DTN in the space setting. One concerns a gateway that terminates the SMTP session locally and passes the message to the Bundle Protocol (7). Another describes how to encapsulate the message into a bundle (8). A third concerns a naming model for off-Earth domains and the need to structure message-authentication components, including DMARC (13), so that only resources local to the network are needed to handle the request (14). BPMail is an open implementation of the gateway approach (9).

Spam is mentioned in just one of those drafts, which states that spam can cause denial-of-service problems for a store-and-forward network and recommends two solutions: checking integrity and authenticity of the sending network, as with DMARC, and whitelisting sending networks (7). These are preventive measures addressing who is allowed to send messages. Neither draft mentions any tool that could examine messages as they arrive. The security section of the encapsulation draft merely notes that a payload carrying bad data is itself a denial of service (8).

2.3 Anti-spam filtering

The spam-detection area is large and focuses mainly on classification accuracy. A survey of machine-learning methods (15) describes major techniques, their comparative advantages, and open issues, in terms of a filter's ability to separate spam from non-spam. A digest-based approach to detecting spam is considered in the same terms: the algorithm underlying collaborative filtering was tested against three criteria, its robustness to automation, robustness to attacks, and false-positive rate (16).

The underlying mechanisms implemented by the tools used in this study are documented by their projects rather than in academic papers. How a blocklist is published and queried as a DNS zone is described by Spamhaus (17).

SpamAssassin's network configuration (18) describes its network rules and the parameters that constrain them. The URI blocklist plugin (19) describes its DNS queries, from resolving name servers to addresses. The sender-authentication techniques SPF, DKIM and DMARC are specified in RFC 7208 (20), RFC 6376 (21) and RFC 7489 (13) correspondingly. rspamd (22) and Vipul's Razor (23) are also documented by their projects.

What this body of work measures is detection rate and false positives. What it does not measure is time. The evaluations are run on a network where a lookup returns in milliseconds, so the cost of reaching a remote service does not appear as a variable.

2.4 DTN security

The DTN security overview (24) defines the security requirements of a store-and-forward network and the rationale for choosing or rejecting specific security measures. The problem it addresses is scarcity of resources on a node and their use by inappropriate traffic. It proposes authenticating the bundle to ensure it

originates from a legitimate entity permitted to use the network. This shares a starting assumption with the present work, that unwanted traffic is the threat, but answers it differently: authentication of the sender rather than classification of the message.

2.5 Related problems

Four problems already studied elsewhere have the same structure as the one examined here.

The first is the round-trip cost embedded in a protocol that is then carried over DTN. The routing study that defined the DTN problem (25) recognised the importance of the DNS query and response, and of a reliable transport exchange, for an SMTP transaction. The SMTP case was handled by terminating the connection at a gateway. It arose again with HTTP encapsulated over the Bundle Protocol: that draft (26) explains that while version negotiation and URL redirection remain possible, they are not recommended, since they require further bundling over the delayed network. The same issue applies to the network lookups a filter performs.

The second is sequential DNS lookups. Spamhaus (27) explains that performing blocklist queries sequentially rather than simultaneously raises the cost of an individual query from tens to hundreds of milliseconds, and that this accumulates across a set of queries, noting that a normal mail transaction involves tens of these queries. This increase is minimal on a terrestrial network, but the principle matters under delay.

The third is that latency of network tests is well recognised by their operators. Advice on the SpamAssassin users' mailing list, archived on the project wiki (28), describes the difference between processor time and latency: network tests, including blocklist and digest tests, cause considerable latency but need little processor time, and it advises allowing up to fifteen or twenty seconds per message so the test does not time out. That figure assumes a network with millisecond lookups.

The fourth one is where the filtering process occurs. One paper relocates the filtering system from the receiving server to the sending server in order to reduce detection time (29).

2.6 Methods and their limits

Each of these fields evaluates its systems in a way that cannot measure behaviour under delay.

Work on DTN applications is evaluated by analysing the protocol specification rather than by running the system. The HTTP encapsulation draft (26) identified the round-trip problem through protocol analysis and designed around it. This works when the round-trips can be seen in the protocol. The network checks a filter makes are not: they happen in the scanner process, driven by plugin behaviour, timeouts, and chains of lookup dependencies that do not appear in any protocol document. The chain of URI blocklist lookups, for example, is specified only in the plugin documentation (19). Protocol analysis cannot find them.

The anti-spam evaluation measures a filter's performance on a corpus over a low-latency network, reporting accuracy (15, 16). Done this way, it produces no timing information, and cannot, because the latency it would measure is not present.

The DTN evaluation uses simulation and emulation, measuring delivery ratio, routing behaviour, and buffer occupancy (25).

This dissertation therefore carries mail over a working DTN stack and runs the tools against it under controlled delay, recording per-message timing, verdicts, and packet captures. Behaviour that is neither written in a specification nor visible in a terrestrial accuracy benchmark can only be found by running the software and observing the network.

2.7 The gap

Email over DTN has been deployed and drafted for, but that work moves mail across a delayed channel without filtering it. The anti-spam literature is concerned with accuracy, meaning how well a classifier separates spam from good mail, and measures it on low-latency networks, so delay never enters the picture. Work on DTN security does treat unwanted traffic as a problem, but its answer is to authenticate senders, not to inspect messages. And the round-trip cost that filtering runs into has been noticed before, in adjacent applications, but it was designed around rather than measured.

The SMTP gateway draft (7) does name spam as a threat to a store-and-forward network, but goes no further than recommending policy. So how the tools actually behave once delay is added to their lookups has not been measured, and that is the question this dissertation sets out to answer.

2.8 Summary

The transport problem for email over DTN has been addressed by deployments and standards work. The filtering problem has not. The filtering tools depend on network services whose cost has only been measured on a low-latency network, and the fields that might study it rely on methods that cannot observe that cost. The next chapter describes the testbed built to measure it directly.

3 | System Design

This chapter describes the testbed used to evaluate anti-spam tools under one-way delay representing a deep-space DTN link. It sets out the node layout, the host model and the reasons for it, the tools under test, the DTN configuration, the delay simulation, and the experimental design.

3.1 System Architecture

The testbed has two networks joined by a relay. Each node runs an ION instance (Section 2.1.2) and is a separate virtual machine running Ubuntu 22.04 LTS. Earth uses 192.168.100.0/24 and Moon uses 192.168.200.0/24. The nodes are:

- **Mail1 (Earth).** The Earth mail server. Runs Postfix and Dovecot. Address 192.168.100.10. ION node 1.
- **Space (Relay).** The transit node between the two networks. Runs an ION node and carries the delay and packet-capture configuration. Earth-side interface 192.168.100.1, Moon-side interface 192.168.200.1. ION node 2.
- **Mail2 (Moon).** The Moon mail server. Runs Postfix, Dovecot, and the three anti-spam tools. Address 192.168.200.11. ION node 3.
- **Client1 (Earth).** User terminal on the Earth network. Runs Thunderbird and a script (the test harness)(Section 4.9) that injects the spam messages. Connects to Mail1 over IMAP (port 143) and SMTP submission (port 587, STARTTLS).

- **Client2 (Moon).** User terminal on the Moon network. Runs Thunderbird and connects to Mail2 the same way.

Mail moves server to server across the DTN link: a message submitted on one mail server is carried as a bundle through the relay to the other. ION runs on all three servers, so the bundle path is node 1 (Earth) to node 2 (relay) to node 3 (Moon). The relay is the only path between the two networks, so all traffic between Earth and Moon passes through it. This is where the one-way link delay is applied, standing in for the Earth–Moon propagation time. Mailboxes use mbox format, and each user reads mail from their local server. The topology is shown in Figure 3.1.




Figure 3.1: Testbed topology. Each client connects only to its local mail server. The relay is the only path between the Earth and Moon networks and is where the one-way link delay is applied.

Data flow (Earth to Moon). The path of a message is shown in Figure 3.2.




Figure 3.2: End-to-end path of an email from Earth to Moon across the DTN link.

3.2 Environment and Technology Selection

3.2.1 Virtual machines over containers

The testbed runs on full virtual machines (Ubuntu 22.04 LTS, ARM64, under UTM on macOS). Containers were considered but not used. They were not used for three reasons:

- **VMs match the topology.** The testbed mimics three separate machines — Earth, Space, Moon — connected by a long-delay link, and full VMs naturally model this, whereas containers share one kernel on one host.
- **Simpler delay setup.** Setting up an intermediate functional block (ifb), the virtual device used to apply delay to incoming traffic (Section 3.7), is fiddly and error-prone inside containers, but it is straightforward on a VM where each node has its own kernel and network stack.

- **DTN state stays isolated per node.** Each ION node keeps its own bundle store and contact schedule, and a full VM ensures the link delay sits on the path itself rather than leaking into a node's own processing.

3.3 Email stack

- **Postfix** was chosen for its clear interface for handing mail to an external program, needed to pass messages to the DTN layer.
- **Dovecot** was chosen because it works directly with mbox format and the local system accounts that Postfix delivers to.
- **Thunderbird** was chosen as a standard mail client, so the message traffic in the testbed matches what a normal user would generate.

3.4 DTN stack

- **ION-DTN** was chosen for its deep-space flight heritage and its existing tool for carrying email over the Bundle Protocol, which connects directly to the mail stack.

3.5 Anti-Spam Tool Selection

Three tools were selected, each with a different external dependency:

- **SpamAssassin 3.4.6** — content scoring that depends on DNS blocklist lookups.
- **rspamd 2.7** — content scoring with a large number of asynchronous DNS lookups.
- **Vipul's Razor 2.84** — collaborative filtering that checks each message against a central catalogue server, which holds signatures of known spam reported by other users.

The three cover the main ways a filter reaches the network: blocklist DNS (SpamAssassin, rspamd) and a TCP request to a central reputation server (Razor). Testing one tool would show only one failure path.

3.6 DTN Design

The DTN layer uses ION-DTN 4.1.4 (30) and BMail. The three servers run ION as nodes on a single ipn scheme (Section 2.1.1): Mail1 (Earth) is node 1, Space (relay) is node 2, and Mail2 (Moon) is node 3. Bundles are carried over a UDP convergence layer (Section 2.1.1), with each node listening on its own port: 1113 on Mail1, 2113 on Space, and 3113 on Mail2. The relay forwards bundles between the two mail nodes, so the path from Earth to Moon is node 1 to node 2 to node 3. The contact plan (Section 2.1.2) sets the windows during which the nodes can exchange bundles. BMail sits above ION on each mail server: it takes a message from Postfix, encapsulates it into a bundle for the destination node, and rebuilds it on arrival. The full configuration, including the SDR settings (Section 2.1.2), the contact plan, and the startup order, is given in the implementation chapter.

3.7 Delay Simulation Design

Delay is applied with netem on the relay, to all traffic crossing the boundary between the Earth and Moon networks. An Earth–Moon link delays travel in both directions: a signal takes about 1.3 seconds each way, so a round-trip takes about 2.6 seconds. The testbed emulates this by delaying each leg on the relay’s Moon-facing interface.

The complication is that tc can only shape traffic *leaving* an interface, so the two legs have to be handled separately. The Earth-to-Moon leg is the egress direction on that interface and is delayed by a netem qdisc attached to the interface root. The Moon-to-Earth leg arrives as ingress, and ingress has no egress queue to attach a

qdisc to. To delay it, incoming traffic is redirected to an Intermediate Functional Block (ifb) device, whose own egress path can be shaped normally, and the same netem delay is applied there. The two mechanisms, the root qdisc and the ifb device are therefore the two halves of a single symmetric delay, not a way to delay one direction while sparing the other. Traffic that stays within a single network never reaches the relay and is unaffected. This is shown in Figure 3.3.




Figure 3.3: Symmetric delay. Each leg crossing the relay is delayed by 1300 ms: the Earth-to-Moon leg on the interface’s egress path, and the Moon-to-Earth leg on an ifb device, since ingress cannot be shaped directly. The two legs together give a round-trip time of about 2.6 s. Traffic that stays within one network does not cross the relay and is not delayed.

Delay is the independent variable. The experiments use a representative one-way delay of 1.3 seconds, a round-trip of about 2.6 seconds, as a deep-space operating point. The testbed allows any value to be set, so the same setup covers the range from short Earth–Moon delays to multi-minute deep-space delays.

3.7.1 DNS Resolver

The blocklist rules in every tool under test resolve DNSBL and URIBL names, so the Moon node needs a working recursive resolver. A large public resolver cannot serve this role: the major blocklist operators refuse or rate-limit queries that arrive through open public resolvers such as those run by Google or Cloudflare, returning an error or an empty answer rather than a listing (17). Scanning through one would report no spam for reasons unrelated to delay and make the results meaningless.

The testbed therefore uses a local resolver, running on an off-testbed virtual machine with ordinary internet access to the live blocklist zones. The Moon node runs a local unbound instance that forwards to it, for the following reasons:

- **Valid blocklist answers.** A dedicated resolver querying the blocklist zones is what those operators expect, so listings are returned normally and all that is added to a lookup is the injected link delay.
- **Control of the resolver's own timers.** A public resolver cannot be reconfigured, yet the timers that decide how a resolver treats a late answer turn out to determine whether detection survives delay (Chapter 5). A private resolver can be tuned and its configuration recorded.
- **Reproducibility.** Its cache and per-server state can be reset before each run and dumped afterwards, so every run starts cold and its DNS state is on record.

The recursive resolver sits off the Moon network, so every query the scanner makes crosses the relay and is delayed like the mail.

3.8 Summary

This chapter set out a five-node testbed for measuring anti-spam tool behaviour under link delay. The design uses full virtual machines for kernel and DNS isolation, a private resolver to give the Moon side a working DNS path, three tools chosen to cover the main external dependencies, and an ifb-based scheme to apply delay. The next chapter describes how this design was built and configured, and the chapter after it presents the results.

4 | Implementation

This chapter describes how the testbed was built and how the experiments are run on it. It covers the DTN and mail layers, the delay mechanism, the resolver, the setup of each anti-spam tool, and the harness that executes a full set of runs without supervision. This chapter is written in enough detail to help future work that builds on this setup.

4.1 Basis and Scope of This Work

The DTN transport layer of the testbed follows the approach established in earlier work on delay-tolerant email, which built and validated an ION and BPMail pipeline for carrying mail over the Bundle Protocol. The same set of software is used, but instead of being containerised, the software stack is built on virtual machines, for the reasons given in the previous chapter.

All other components of the testbed architecture are new: delay introduction on the relay, the external resolver, installation and setup of all anti-spam tools, the measurement methodology, and the experimental design.

4.2 Environment and Dependencies

The testbed uses five VMs on an Apple Silicon MacBook Air, under UTM. All five run Ubuntu 22.04 LTS (ARM64). Three servers are named mail1, space and mail2; two clients, one for each network. The mail, DTN and anti-spam applications are

installed on the servers. Clients run Thunderbird and serve only to verify that mail transfer works end-to-end; they do not participate in measurement.

Three libraries are more recent than those in the Ubuntu 22.04 repositories and must be built from source on each server before ION and BPMail are compiled: mbedtls 2.28.8, gmime 3.2.7 and c-ares 1.34.4. Each is built in order, then `ldconfig` is run so the libraries are found first.

4.3 DTN Layer

ION-DTN 4.1.4 is built from source on all three servers:

```
./configure --disable-crypto CFLAGS="-Dlinux"  
make  
sudo make install  
sudo ldconfig
```

The `-Dlinux` flag is required because ION's `platform.h` selects the system headers that define `MAXPATHLEN` behind an `#ifdef linux` check. Crypto is disabled because the Bundle Protocol Security extensions are not used here.

Each node is configured from the `bench-udp` demo template supplied with ION. The node numbers are 1 for mail1, 2 for space and 3 for mail2. Bundles are carried over the UDP convergence layer, with each node listening on its own port: 1113 on mail1, 2113 on space and 3113 on mail2. The bundle protocol configuration for mail1 is:

```
a scheme ipn 'ipnfw' 'ipnadminep'  
a endpoint ipn:1.0 x  
a endpoint ipn:1.1 x  
a endpoint ipn:1.128 x  
a endpoint ipn:1.129 x  
a protocol udp 1400 100  
a induct udp 192.168.100.10:1113 udpcli
```

```
a outduct udp 192.168.100.1:2113 udpclo
```

Space is the only node with two inducts and two outducts, and is the only route between the two networks. Routing uses Contact Graph Routing, so routes are computed from the contact plan rather than configured statically. The contact plan is identical on all three nodes and declares reachability in both directions between each adjacent pair:

```
m horizon +0
a range +0 +86400 1 2 1
a range +0 +86400 2 1 1
a range +0 +86400 2 3 1
a range +0 +86400 3 2 1
a contact +0 +86400 1 2 100000000
a contact +0 +86400 2 1 100000000
a contact +0 +86400 2 3 100000000
a contact +0 +86400 3 2 100000000
```

This contact plan is simple and more-or-less always on: the contact windows are set to 24 hours, since shorter windows expire during a run and bundles stop being forwarded, so the window has to cover a full pass of three runs with margin. A more complex scenario, such as one including an orbiter whose visibility to a ground station comes and goes, would need a contact plan with scheduled outages instead, but that is not needed to study the effect of delay alone.

4.4 Mail Layer

BPMail 0.1.0 provides the gateway between the mail system and ION. It is installed on mail1 and mail2 only; space forwards the bundles between them without running BPMail or speaking any mail protocol, even though those bundles carry mail as their content.

4.4.1 Build

BPMail is built with Meson. One change to the build configuration is required. The first build failed with an error reporting that no value was defined for MAXPATHLEN. The cause is an interaction between two settings. ION's `platform.h` includes the system headers that define MAXPATHLEN only when the `linux` platform macro is set, guarded by `#ifdef linux`. BPMail's `meson.build` specified `c_std=c17`, and strict ISO C17 disables non-standard platform macros such as `linux`, so the guard did not pass and the definition was never picked up, even when `-Dlinux` was passed on the command line. Changing the standard to GNU C17 keeps those platform macros while remaining C17-compliant:

```
sed -i "s/c_std=c17/c_std=gnu17/" ~/bpmail/meson.build
```

The platform flag is supplied through a native file, and the build then completes:

```
printf '[built-in options]\nc_args = [\x27-Dlinux\x27]\n' > linux-native.txt
meson setup build --native-file linux-native.txt
ninja -C build
sudo ninja -C build install
```

4.4.2 DTPC configuration

BPMail transfers mail over DTPC (Section 2.1.2) rather than raw bundles, so the two mail nodes need DTPC endpoints and a profile. Service 128 is used for sending and 129 for receiving, and the endpoints are added to each node's bundle protocol configuration as shown above. The DTPC profile is the same on both nodes:

```
1
w 1
a profile 1 5 1000 10 600 0.1 dtn:none
s
```

The profile sets a maximum of five retransmissions, an aggregation size limit of 1000 bytes, an aggregation time limit of 10 seconds and a bundle lifetime of 600 seconds. The `dtpcadmin` step is appended to the `ionstart` script on both mail nodes so that DTPC comes up with the rest of ION.

A DTPC topic can only be opened once per node, for sending or receiving but not both, so each direction uses its own topic. The experiments send from mail1 to mail2, and use topic 26 for that direction. The destination endpoint is the receiving service on node 3, `ipn:3.129`; sending to the `service-128` endpoint does not reach the receiver.

4.4.3 Postfix and Dovecot

Postfix is configured on both mail nodes with a transport map that routes mail for the remote domain to a pipe transport, which hands the message to BPMail. BPMail needs access to ION's shared memory, which depends on the working directory, so the pipe calls a small wrapper that changes to the ION node directory first:

```
bpmail unix - n n - 1 pipe
    flags=F user=<user> argv=/home/<user>/bpmailsend_wrapper.sh 1 ipn:3.129
```

A background daemon runs `bpmailrecv` in a loop on each node and pipes what it receives into `sendmail`, which delivers it to the local mailbox. Dovecot serves those mailboxes over IMAP, in mbox format.

This path is what the Thunderbird clients use, and how the testbed was verified as a working mail system. It is not the path used for measurement. The experiments call `bpmailsend` and `bpmailrecv` directly, for the reasons given in Section 4.9.4.

4.5 Delay Simulation

The link delay is applied on space with `tc` and the `netem` queueing discipline, on the Moon-facing interface, so it acts on all traffic crossing the boundary between the two

networks. As explained in 3.7, tc shapes only egress, so the outbound leg uses a netem qdisc on the interface root and the inbound leg is redirected to an ifb device carrying the same delay. The full sequence is:

```
sudo modprobe ifb
sudo ip link add ifb0 type ifb
sudo ip link set ifb0 up
sudo tc qdisc add dev enp0s2 root netem delay 1300ms
sudo tc qdisc add dev enp0s2 handle ffff: ingress
sudo tc filter add dev enp0s2 parent ffff: protocol ip u32 match u32 0 0 \
    action mirred egress redirect dev ifb0
sudo tc qdisc add dev ifb0 root netem delay 1300ms
```

Clearing the delay removes both qdiscs, the ingress filter and the ifb device.

The one-way delay is a parameter. The experiments use 1300 ms in each direction, giving a round-trip time of about 2.6 s. Because the delay sits at the boundary of the Moon network, everything mail2 exchanges with anything outside it pays that cost: the bundles it receives from mail1, and the DNS queries the scanner makes while examining a message.

4.6 DNS Resolver

The scanners run on mail2 and need working DNS. The queries are served by a local unbound instance on mail2, which forwards to a private recursive resolver running on an Azure virtual machine on port 5353.

Two changes to the default unbound configuration were needed.

The first was a startup failure. With DNSSEC validation enabled and all queries forwarded, unbound logged repeated failures to prime the trust anchor and refused to resolve anything. Forwarding and validation conflict in this configuration. Because the upstream resolver performs validation, the validator module can be removed on

mail2, leaving only the iterator:

```
module-config: "iterator"
```

The second concerns two timers that interact with the link delay, both given in milliseconds. For the delayed runs both were increased above unbound's defaults, so that they exceed the round-trip time of the delayed link; the reason these values matter, and what happens with the defaults in place, is a result of the experiments and is presented in Chapter 5.1.1.

```
discard-timeout: 9000 //default: 1900  
infra-cache-min-rtt: 4000 //default: 50
```

4.7 Anti-Spam Tools

Three tools are installed on mail2, the node that receives and scans. Each is invoked once per message, on the message file, and each is run under two configurations: a default one, and one with its timeout reduced below the round-trip time of the link.

4.7.1 SpamAssassin

SpamAssassin is invoked through its standalone binary rather than through

spamc:

```
spamassassin < "$FILE"
```

The client and daemon pair was tried first and had to be abandoned. When spamc cannot reach spamd it fails open, passing the message through unchanged, so every message came back with no X-Spam-Status header and no indication that anything had gone wrong. The standalone binary scans without a daemon, and re-reads `local.cf` on every invocation, which means a configuration change takes effect on the next message with no restart.

The two configurations differ only in whether the timeouts are set. The default profile removes any timeout lines from `local.cf`, leaving SpamAssassin's own defaults of 15 s for `rbl_timeout` and 5 s for `spf_timeout`. The reduced profile sets both to 1.5 s, below the 2.6 s round-trip time:

```
sudo sed -i '/^rbl_timeout /d;/^spf_timeout /d' /etc/spamassassin/local.cf
echo 'rbl_timeout 1.5' | sudo tee -a /etc/spamassassin/local.cf
echo 'spf_timeout 1.5' | sudo tee -a /etc/spamassassin/local.cf
```

4.7.2 rspamd

`rspamd` runs as a daemon and is invoked through its client:

```
rspamc --ip 1.2.3.4 --mime < "$FILE"
```

The `-ip` option supplies a client address so the message is treated as having arrived from outside. Without it, mail presented locally is treated as internal and the network checks central to this study are skipped.

`rspamd` caches state in the running daemon, so it is restarted before each run to start from a cold cache. The reduced-timeout profile writes a single override file and restarts:

```
echo 'task_timeout = 1.5s;' | sudo tee /etc/rspamd/local.d/options_override.inc
sudo systemctl restart rspamd
```

The default profile removes that file and restarts.

4.7.3 Vipul's Razor

Razor is invoked with `razor-check`, which reports its verdict through the exit code: 0 if the message is catalogued as spam, 1 if it is not, and 2 on a network error. The third case matters here, because it separates a decision from a failure to reach the server:

```
razor-check < "$FILE"
```

Razor's server list is refreshed with `razor-admin -discover` before each run.

Razor differs from the other two in one respect that affects the experiment directly: it has no timeout setting in `razor-agent.conf`. Its connection timeouts are fixed in the client source. Producing the reduced-timeout condition therefore requires patching the module that holds them, keeping a copy of the original so the default condition can be restored:

```
sudo cp -n "$CORE_PM" "$CORE_PM.orig"
sudo sed -i 's/Timeout *=> *20/Timeout => 2/g' "$CORE_PM"
```

That a source patch is the only way to change this value is itself a finding about the tool, and is returned to in Chapter 5.

4.8 Email Corpus

The corpus is drawn from a set of spam messages provided by Richard Clayton, captured from live mail traffic in June 2026. 681 of these were used in every run. Because spam content and the blocklists used to catch it both change over time, this corpus reflects spam as it appeared at that point and is not expected to remain representative, so it is unlikely to be reusable for a repeat of this experiment in future without being refreshed; blocklist listings alone were observed to drift measurably over a two-day gap between runs in this study (Section 5.1.3).

4.9 Run Harness

The three runs for one tool take several hours and involve three machines, so the experiments are automated. The harness is a Python program that runs on the Mac and drives the testbed over SSH. A single command performs a complete pass:

```
python3 run.py --tool spamassassin
```

Every parameter it uses is held in one configuration file, so the procedure is identical between tools and between passes, and the file itself records the exact settings under which a set of results was produced.

4.9.1 Structure

The harness connects to the three servers with Paramiko, a Python library for scripting SSH sessions. Mail1 and mail2 are not reachable directly from the Mac, so their connections are tunnelled through space, which is opened first.

Work on the servers is done by two shell agents, one on mail1 that sends and one on mail2 that receives and scans. They are written out to each node at the start of a run, together with an environment file holding that run's settings, and launched detached so that they survive the SSH session. Each agent reports through two files in its working directory: STATUS, which holds RUNNING, DONE or FAILED, and progress, which holds a count of messages handled. The host polls these every thirty seconds and prints a line for each. The alternative, holding a live channel open to each node for the length of a run, is fragile over hours; sentinel files survive a dropped connection.

4.9.2 Preflight and boot

The harness controls the virtual machines through utmctl, UTM's command-line tool. Before anything starts, it checks that utmctl is usable, that the named virtual machines exist, and that the results directory has at least 20 GB free. It then starts the machines in order and waits for each to accept SSH, up to five minutes. The harness also runs caffeinate, a macOS utility that prevents the machine from sleeping, for the length of the session so the Mac does not go to sleep partway through a run.

4.9.3 Checks before a run

Before a run starts, the harness checks that the testbed is actually in a state where the results will mean something. Any of the checks below, called gates, can stop the run before it begins.

- **Clocks.** The sender on mail1 and the receiver on mail2 need to agree on the time, since their timestamps are compared to measure delay. The harness checks each node's clock against the host and stops the run if either is off by more than a second.
- **Resolver.** unbound has to be running on mail2 and able to answer a test query, or none of the DNS-dependent tools would work.
- **ION.** ION is restarted on all three nodes. This also resets the contact window so it covers the whole run.
- **Tool.** The harness asks the tool for its version, and a response confirms it is installed. That version string is saved alongside the results.

4.9.4 Sequence of a run

Each run follows the same sequence.

The testbed is first returned to a cold state: unbound is restarted, which clears both its cache and its record of upstream server behaviour, any agent state from a previous run is deleted, the tool is reset, and any leftover netem configuration is removed. The resolver's infrastructure cache is dumped and saved, so that its state before the run is on record.

The tool profile for the run is then applied, along with the delay if the run calls for it. The round-trip time is measured by pinging mail1 from mail2 and checked against a window: 2400 to 2900 ms for a delayed run, and under 150 ms for the baseline. A measurement outside the window aborts the run. This is the most important gate in

the harness. A run whose delay was not actually in place, or was applied twice, would produce plausible-looking numbers that mean nothing, and the failure would not be obvious afterwards.

Packet capture is then started on space and on mail2, excluding management SSH so that the harness's own traffic does not appear in the bandwidth figures:

```
sudo tcpdump -i any -s 0 -w {out} -C 100 -W 30 'not tcp port 22'
```

The receiver is launched first and the sender three seconds later, so that no message can arrive before the receiver is listening.

The sender goes through the corpus in sorted order. For each message it records the Message-ID and a timestamp, sends it, and records the exit code:

```
mid=$(grep -i -m1 '^Message-ID:' "$f" | tr -d '\r\n' | cut -d' ' -f2-)
sent=$(date -u +%s.%N)
eval "$SEND_CMD"
rc=$?
echo "$seq,$bn,$mid,$sent,$rc" >> "$SENT"
```

The send command pipes the message directly to BMail:

```
sed "1{/^From /d}" "$FILE" | bmailsend -t 26 1 ipn:3.129
```

The leading From line of the mbox-format corpus file is stripped, because it is not part of the message.

The receiver blocks on bmailrecv until a bundle arrives, writes the message out, scans it, and records the timing:

```
RAW=$(timeout "$RECV_TIMEOUT" bash -c "$RECV_CMD")
arrived=$(date -u +%s.%N)
printf '%s' "$RAW" > "$eml"
start=$(date -u +%s.%N)
eval "$SCAN_CMD" > "$outf" 2>&1
```

```
rc=$?  
end=$(date -u +%s.%N)
```

The scan output is saved in full for every message, not just the parsed verdict, so that the parsing can be revised later without repeating a run.

Sending and receiving through BPMail directly, rather than through Postfix, gives a clean boundary for the scan time. The clock starts when the message arrives and stops when the tool exits, so the recorded time is the time the tool spent on that message. Invoking the tool through Postfix would include its queueing in every measurement, and at a round-trip time of 2.6 s the two could not be separated.

Messages are paced rather than sent back to back. The interval differs by run and is set in the configuration.

4.9.5 Completion and collection

When the sender reports DONE the harness enters a drain phase, because messages are still in flight and will keep arriving for at least one one-way delay afterwards. It waits for the receiver's count to reach the number sent, up to fifteen minutes, then writes a STOP file that the receiver checks between bundles, so it exits cleanly rather than being killed mid-scan.

Captures are stopped, and the resolver's state is recorded again. Two health figures are taken from it: the highest retransmit timeout in the infrastructure cache, and the number of log entries reporting answers discarded for arriving too late. A run whose highest retransmit timeout has reached 120,000 ms is marked SUSPECT in its manifest, because that is the ceiling and indicates the resolver stopped working properly during the run. The check runs automatically, so a compromised run is flagged at the time rather than discovered later.

All artifacts are then pulled back to the Mac. The sent and received records are reconciled by Message-ID to identify any messages that did not arrive, and the

receiver's timing records are joined with the parsed scan output into a single per-message file. Each tool has its own parser, because the three report their verdicts differently: SpamAssassin in an X-Spam-Status header that folds across lines, rspamd in X-Spam-* headers added to the message, and Razor in its exit code alone.

Each run finally writes a manifest recording the tool and version, the profile, the measured round-trip time and the gate against which it was checked, the timings, the message counts, the resolver health figures, and the resulting status. The manifest is what makes a set of results interpretable months later.

4.10 Running an Experiment

A pass consists of three runs against the same corpus of 681 messages:

- **Run 1**, no delay, default tool configuration. The baseline.
- **Run 2**, delay applied, default tool configuration.
- **Run 3**, delay applied, tool timeout reduced below the round-trip time.

A completed pass leaves, for each run, the sender's record, the receiver's timing record, every scan output, packet captures from space and mail2, the resolver state before and after, the joined per-message file, and the manifest. A summary across the three runs is written at the end of the pass.

4.11 Summary

The testbed is a working DTN mail system with a controllable link delay, a resolver that reaches live blocklists, and three anti-spam tools, each configurable into a default and a reduced-timeout condition. The experiments are driven by a configuration-file-based harness that verifies the testbed before each run, refuses to proceed if the delay is not correct, collects every artifact, and flags runs whose

resolver state indicates the results cannot be trusted.

5 | Results and Discussion

5.1 SpamAssassin

SpamAssassin was run three times over the 681-email corpus: a no-delay baseline, Run 2 (delay, default timeouts), and Run 3 (delay, reduced timeouts). Each message was scanned once by the standalone binary and timed from arrival to exit.

Metric	Baseline	Run 2	Run 3
Flagged YES	383 (56.2%)	355 (52.1%)	316 (46.4%)
Emails with DNSBL hit	416	400	251
Scan min/avg/max (s)	0.45/1.68/7.58	3.17/7.36/16.50	0.43/4.05/14.83

Table 5.1: SpamAssassin - Aggregate results across the three runs.

5.1.1 Detection

With no delay, 383 messages were flagged. Applying 1.3 s each way then broke detection at the resolver: DNSBL answers returned at about 2.64 s (one RTT). But, they were discarded by unbound before reaching the scanner with the error message drop reply, it is older than discard-timeout. The default value of discard-timeout, 1900 ms, is lower than the RTT of 2600 ms, thus all delayed responses were too old and discarded (31). Setting discard-timeout to 9000 and infra-cache-min-rtt to 4000, which were higher than the RTT, solved the problem, and a cold resolution took 2.694 s while the test message regained its score.

With the resolver fixed, Run 2 nearly preserved detection (416 → 400 DNSBL emails) at about four times the scan time. Run 3 cut rbl_timeout and spf_timeout to 1.5 s,

below the RTT, trading detection for speed. Two things follow. First, every DNSBL rule fired less than in Run 2 and none fired more (Table 5.2).

Rule	Baseline	Run 2	Run 3
RCVD_IN_PBL	151	149	73
RCVD_IN_SBL_CSS	148	136	43
URIBL_DBL_SPAM	118	117	73
RCVD_IN_PSBL	108	93	67
RCVD_IN_MSPIKE_BL	103	93	32
RCVD_IN_XBL	82	75	20
RCVD_IN_BL_SPAMCOP_NET	75	63	12

Table 5.2: SpamAssassin - DNSBL/URIBL rule firing (selected rules); every rule falls from Run 2 to Run 3.

Each rule corresponds to one blocklist: RCVD_IN_* rules check the sending server’s IP against an address-based list, and URIBL_* rules check domains in the message body against a domain-based list. (Table 5.2) shows the individual DNSBL/URIBL rules; every rule falls from Run 2 to Run 3.

Second, the resolver failed after approximately email 500. With `rb1_timeout` and `spf_timeout` reduced to 1.5 s, SpamAssassin abandoned DNSBL lookups before the replies could return across the 2.6 s DTN round-trip. However, the delayed replies still reached unbound. From the resolver’s perspective, the upstream DNS server had failed to respond before its retransmission timeout (RTO), so the query was treated as an upstream timeout and the RTO for that server was increased using exponential backoff. As this occurred repeatedly across hundreds of emails, the RTO eventually reached Unbound’s maximum value of 120 000 ms, as defined by the default `infra-cache-max-rtt` setting (31).

According to the Unbound documentation (32), reaching this value places the upstream server into the *blocking regime*, where Unbound considers the upstream server unresponsive and temporarily stops sending DNS queries to it and Unbound will attempt to use the server again only after the infrastructure-cache entry expires. Because the experiment used a single upstream resolver, subsequent DNSBL lookups could no longer be resolved successfully. Consequently, after approximately

email 500, SpamAssassin effectively scanned messages without DNS-based checks, causing the DNSBL hit rate to collapse (Table 5.3) and scan times to fall sharply. The resolver remained in this blocked state until it was restarted.

Corpus position	Run 2	Run 3
Emails 1–500	260/500 (52%)	220/500 (44%)
Emails 501–681	140/181 (77%)	31/181 (17%)

Table 5.3: SpamAssassin - DNSBL-hit rate by corpus position; Run 3 collapses after $\approx$ email 500.

The 251 DNSBL-hit emails in Run 3, despite a 1.5 s timeout below the 2.6 s RTT, are explained by the adaptive nature of `rbl_timeout`. It is not a hard per-query cutoff but the starting value of a timeout that shrinks as queries complete, with a one-second grace floor granted to each still-pending query. Messages whose lookups form a serial dependency chain — the URIBL domain $\rightarrow$ NS $\rightarrow$ IP $\rightarrow$ DNSBL wave chain, or nested SPF include: directives — launch fresh pending queries at each new wave, and the grace floor keeps extending the effective deadline past the RTT, so the delayed answers arrive in time to score. The correlation confirms this: emails that scanned longer than the RTT hit a DNSBL rule 47% of the time, against 24% for emails that scanned faster, and 83% of all Run 3 hits came from the longer-scanning group. Cache reuse of shared-domain listings is a secondary contributor. The result is that the same serial dependency chains that increase scan time under network delay also preserve DNSBL detection when the timeout is reduced, because the longer scan duration allows delayed DNS replies to arrive before scanning completes.

5.1.2 Scan time

Even in baseline conditions, the scan is not instantaneous, the network rules depend on the DNSBL replies, which can take time, and the scan concludes only when the last one of them arrives; with the resolver set up properly, delayed scans were still taking around 8 seconds. Since the lookups go in serial waves, dependent on the results from the previous waves, this suggests that the latency depends on the depth

of the dependency chain, not the number of lookups. There are two causes of such depth – URIBL chain, which resolves the name servers of the domain into IPs and performs DNSBL lookup of those IPs (domain > NS > IP > DNSBL), and nested SPF include: directives, each fetched before its children are visible.

Run	<1 s	1–3 s	3–6 s	6–10 s	>10 s
Baseline	438	101	104	38	0
Run 2	0	0	316	213	152
Run 3	162	129	194	184	12

Table 5.4: SpamAssassin - Scan-time distribution (emails per bucket).

The distribution changes completely with delay, since no Run 2 message scan completes in less than 3 s, since all must now do at least one round-trip time and dependency-chain messages must do many. Run 3 splits into two groups: the shorter timeout cuts many scans short, while dependency-chain messages still run long.

5.1.3 Flips

	Run 2 → Run 3
Verdict unchanged	424
Flipped YES → NO	148
Flipped NO → YES	109
Net flagged change	−39
Gained a DNSBL rule (drift)	178 (26.1%)

Table 5.5: SpamAssassin - Verdict changes and blocklist drift, Run 2 to Run 3.

Run 3 was done two days later than Run 2, so the effect for this pair is a combination of both. The shortened timeout aborts the DNS query before the answer comes back, eliminating DNSBL matches and shifting messages from YES to NO. However, 178 messages (26.1%) *acquired* the DNSBL rule between the runs, and a shortened timeout cannot do this, so this is blocklist drift over the two days. The 109 NO → YES switches are thus drift, while the 148 YES → NO switches are a mixture of timeout failure and drift.

5.1.4 Network volume

Component	Baseline	Run 2	Run 3
DNS bytes	5,160,681	5,418,742	3,413,349
DNS packets	37,789	39,875	25,783
Email Bundle bytes	572,625	538,529	547,709
Email Bundle packets	1,447	1,369	1,389
DNS % of bytes	90.0%	91.0%	86.2%

Table 5.6: SpamAssassin - Traffic on the relay: DNS to the resolver vs DTN bundle traffic.

Traffic in DNS is nine to ten times greater than the DTN bundle traffic carrying the email messages in the baseline and Run 2, while the amount of email bundle traffic stays almost the same across all runs. In Run 3, the lower DNS traffic indicates a resolver failure rather than a bandwidth saving. After about email 500, the resolver stopped sending DNS queries upstream, so fewer DNS packets were generated and SpamAssassin could no longer perform DNSBL lookups.

5.2 rspamd

This section reports the three runs with `rspamd`. The runs follow the same corpus and DNS path used by the previous tool. `rspamd` was executed on the receiver using `rspamc -ip 1.2.3.4 -mime`; the `-ip` option ensures that the external processing is done, and the rules dependent on DNS run. Runs 1 and 2 were executed using the default task timeout, which was 8 s, while Run 3 uses 1.5 s, below the round-trip time.

5.2.1 Detection

Metric	Run 1 (baseline)	Run 2 (delay)	Run 3 (delay, 1.5 s)
Messages processed	681	681	681
Flagged as spam	675 (99.1%)	654 (96.0%)	569 (83.6%)
Messages with real hit	502	482	94
Messages with _FAIL	0	45	665
Total real hits	1449	1330	221
Total _FAIL	0	357	2729
Average score	18.13	17.20	9.26
Maximum score	46.10	45.55	21.55

Table 5.7: rspamd - Detection and scoring across the three runs.

Baseline flags 675 of 681 messages (99.1%); this is the detection ceiling. Under delay with default timeouts (Run 2) detection is nearly preserved at 654 (96.0%). Reducing `task_timeout` to 1.5 s (Run 3) drops it to 569 (83.6%). The cause is in the hit and _FAIL rows: real hits fall 502 → 482 → 94 as timed-out lookups rise 0 → 45 → 665.

5.2.2 Scan time

Scan time (s)	Run 1	Run 2	Run 3
<1 s	644	0	2
1–2 s	24	0	679
2–4 s	1	113	0
4–6 s	6	428	0
≥6 s	6	140	0

Table 5.8: rspamd - Scan-time distribution across the three runs (messages per bucket).

Baseline scans fast: 644 out of 681 scans in less than 1 s. In the delay setup with default timeouts (Run 2) every scan is slow (median 5.67 s), where the round trip time dominates since every uncached lookup incurs one RTT cost; 140 out of 681 scans hit the ceiling at 8 s. In Run 3, scan time reduced to 1 to 2 s for 679 out of 681 scans, with none above 1.64 s.

5.2.3 Flips, Run 2 to Run 3

	Run 2 → Run 3
Verdict unchanged	560
Flipped YES → NO	103
Flipped NO → YES	18
Net flagged change	−85
Messages that lost score	579
Mean score change	−7.94

Table 5.9: rspamd - Per-message flips between the two delayed runs.

Reducing `task_timeout` flipped 103 messages from spam to not-spam and lowered the score of 579 of 681. The 18 messages that gained detection run counter to a shorter timeout and are most likely blocklist-listing drift between the two runs rather than a timeout effect.

5.2.4 Network volume

Metric (receiver capture)	Run 1	Run 2	Run 3
DNS queries sent	38,873	80,890	24,116
DNS responses received	38,771	80,822	24,078
Queries to upstream	13,894	19,033	5,969
Total packets	89 k	182 k	59 k
Total data	18 MB	32 MB	14 MB

Table 5.10: rspamd - Network volume per run.

The same number of sent and received queries is seen for all runs, meaning that there were no lost queries on the wire, while changes in traffic volume are due to rspamd itself, not to packet loss.

Run 2 sent almost twice the number of queries than the baseline ($38,873 \rightarrow 80,890$), and there was a proportional increase in data size ($18 \rightarrow 32$ MB). The reason is retransmission. `rspamd` retries sending the DNS request every time it times out waiting for the response, with the number of retransmits equal to the retransmit count, which gives a total per-request timeout of $\text{timeout} \times \text{retransmits}$, which is 1 s

· 5 by default. As the timeout per attempt is shorter than the 2.6 s round trip time, the unanswered queries will be resent until they get their response. So, an entirely fresh lookup will be sent 3 times on average until the answer is received at 2.6 s.

Run 3 does the opposite: query volume falls below the baseline (80,890 $\rightarrow$ 24,116 queries, 32 $\rightarrow$ 14 MB). Since the timeout is only 1.5 seconds, the task is terminated before the DNS budget expires (it is 5 seconds). Therefore, a request can be issued only twice before termination, stopping the sequence of retransmissions that inflated Run 2. This lower volume is the network-side view of the 2,729 _FAIL symbols: the queries went out but were abandoned before their answers returned.

5.3 Vipul's Razor

This section reports the three runs with Razor. The runs follow the same corpus and DNS path used by the previous tools. Razor 2.84 was executed on the receiver using `razor-check < message.eml`. Razor exposes no network timeout through `razor-agent.conf`, so the connect timeout is a hardcoded constant in `Razor2::Client::Core`. Runs 1 and 2 used the default of 20 s, while Run 3 patched it to 2 s, below the round-trip time.

5.3.1 Detection

Metric	Run 1 (baseline)	Run 2 (delay)	Run 3 (delay, patched)
Messages processed	681	681	681
YES (exit 0, spam)	22 (3.23%)	19 (2.79%)	0 (0%)
NO (exit 1, not spam)	659 (96.77%)	656 (96.33%)	0 (0%)
FAIL (exit 2, error)	0 (0%)	6 (0.88%)	681 (100%)

Table 5.11: Razor - Verdict distribution across the three runs.

Under delay with default timeouts (Run 2), the verdict remains essentially the same (3.23% $\rightarrow$ 2.79% flagged). Run 3 collapses the entire detection process as it reports exit 2 in every scan performed. In contrast to the DNSBL-based tools, Razor has only

one cooperative detection method, which means that if the connection fails, it reports no verdict at all.

5.3.2 Scan time

Scan time (s)	Run 1	Run 2	Run 3
<i>distribution (messages)</i>			
<1 s	676	0	0
1–5 s	4	0	0
5–10 s	0	90	671
10–15 s	0	465	10
15–30 s	1	73	0
30–60 s	0	44	0
≥60 s	0	9	0

Table 5.12: Razor - Scan-time distribution across the three runs.

Baseline scans are completed within one second (676 of 681). With delay without timeout (Run 2), the median increases to 11 s, and the distribution gets a heavy tail up to 122.6 s. With the timeout (Run 3), the time of scanning becomes deterministic with a value close to 8 s, which means four connection attempts of the server, each taking 2 s.

5.3.3 Flips

Transition (Run 1 → Run 2)	Messages	Transition (Run 2 → Run 3)	Messages
NO → NO (unchanged)	653	NO → FAIL	656
YES → YES (unchanged)	19	YES → FAIL	19
NO → FAIL	6	FAIL → FAIL	6
YES → NO	3		

Table 5.13: Razor - Per-message flips, matched on a Message-ID and Delivery-date key.

Under delay alone the verdict is stable: 672 of 681 messages keep their baseline verdict, with 3 flipping YES→NO and 6 returning a network error. The timeout run flips 675 of 681 from a valid verdict to FAIL.

5.3.4 Network volume

Metric (receiver capture)	Run 1	Run 2	Run 3
Frames captured	12,814	25,997	57,937
Bytes captured	3.4 MB	5.1 MB	40.2 MB
TCP connections to :2703	687	887	5,460
Connections per message	1.01	1.30	8.02
DNS queries / responses	36 / 42	650 / 810	777 / 971

Table 5.14: Razor - Network volume per run.

DNS traffic is less than 1% of volume in all runs: Razor’s per-check DNS traffic is small, and its delay sensitivity is in TCP, not DNS. That is the opposite of rspamd, where DNS traffic dominated the delayed run.

Connections per message increase from $1.01 \rightarrow 1.30 \rightarrow 8.02$. In the baseline, each message establishes, on average, one TCP connection to a Razor server, indicating that the first selected server usually responds successfully. With delay, some scans use additional servers due to failure; with the timeout, every handshake fails, so nextserver walks through every server in the catalogue and nomination lists until they are exhausted, and all of them are tried on each message. Therefore, the volume of Run 3 (57,937 frames, 40.2 MB) is the largest despite the shortest wall-time: every handshake failure has a kernel-level SYN retransmission, which accounts for roughly two SYNs per attempt across four servers in 681 messages.

5.4 Summary

Tool	Run	Observed	Reason
SpamAssassin	Run 2 (delay, default)	Detection nearly preserved; scan time about four times higher	Tuned resolver timers above the RTT let delayed DNSBL answers arrive; scans block waiting one RTT per lookup wave
	Run 3 (delay, reduced)	Per-email detection falls; resolver collapses near email 500	Timeout below RTT abandons queries early; late answers drive unbound's rto to its ceiling, which does not self-recover
rspamd	Run 2 (delay, default)	Detection nearly preserved; DNS traffic roughly doubles	Answers arrive within the 8 s ceiling; 1 s per-attempt DNS timeout below the RTT retransmits each uncached query
	Run 3 (delay, 1.5 s)	Detection drops; hits replaced by _FAIL; DNS volume falls	Task aborts before answers return, so lookups time out instead of listing; retransmit sequence is cut short
Razor	Run 2 (delay, default)	Detection essentially unchanged; scan time rises with a long tail	Verdicts still returned over the link; slow servers hit the hardcoded read timeout, adding failover delay
	Run 3 (delay, patched)	Detection collapses completely; every check returns FAIL	2 s connect timeout is below the 2.6 s RTT, so every handshake aborts before its SYN-ACK returns
All tools	Delay	Filter DNS/TCP traffic dominates the link, far above the mail	Network rules issue many lookups per message; the mail payload is small by comparison

Table 5.15: Summary of tool behaviour under delay.

Tool	Configuration	Impact under delay
SpamAssassin	<code>rbl_timeout</code> Default: 15 s Run 3: 1.5 s	Below the 2.6 s RTT, so DNSBL lookups are abandoned before their answers return and per-message detection falls.
	<code>spf_timeout</code> Default: 5 s Run 3: 1.5 s	Below the RTT, so SPF lookups do not resolve within the scan.
rspamd	<code>dns.timeout</code> Default: 1 s Tested: 1 s	Sets the per-attempt DNS wait; the effective wait is <code>timeout × retransmits</code> .
	<code>dns.retransmits</code> Default: 5 Tested: 1	Reducing the effective DNS wait below the RTT does not drop detection: a late answer arriving while the task still runs is still used.
	<code>task_timeout</code> Default: 8 s Run 3: 1.5 s	The whole-message ceiling, below the RTT, aborts the scan before answers return; blocklist hits become <code>_FAIL</code> timeouts and DNS volume falls.
Razor	Connect Timeout Default: 20 s Run 3: 2 s	Below the RTT, so every catalogue-server handshake aborts before its SYN-ACK returns; detection collapses to 100% network error.
	<code>can_read</code> Default: 15 s Unchanged	Governs the wait for a server's reply; a server that accepts but is slow to respond stalls the scan up to 15 s, producing Run 2's long scan-time tail.
unbound (resolver)	<code>discard-timeout</code> Default: 1900 ms Configured: 9000 ms	Must exceed the RTT, otherwise delayed DNSBL answers are dropped as too old and detection fails silently.
	<code>infra-cache-min-rtt</code> Default: 50 ms Configured: 4000 ms	Must exceed the RTT, otherwise the resolver treats a delayed answer as a lost packet on first contact.

Table 5.16: Timeout settings that govern behaviour under delay, and the effect of setting them relative to the round-trip time (RTT $\approx$ 2.6 s).

Tool	Baseline	Run 2	Run 3
SpamAssassin	5.16 MB DNS	5.42 MB DNS	3.41 MB DNS
rspamd	18 MB total	32 MB total	14 MB total
Razor	3.4 MB total	5.1 MB total	40.2 MB total

Table 5.17: Summary of network traffic volume across all tools. SpamAssassin reports DNS traffic only, rspamd reports total captured traffic, and Razor reports total captured traffic at the receiver.

SpamAssassin generated approximately 5 MB of DNS traffic in the baseline and Run 2, with traffic falling in Run 3 because the resolver failed after approximately 500 messages. rspamd produced the highest traffic in Run 2 (32 MB) because delayed DNS replies triggered repeated retransmissions. Razor’s network volume peaks in Run 3 (40.2,MB) because repeated TCP connection failures cause nextserver to contact every available server, generating many SYN retransmissions.

Observed (common to all three tools)	Reason
Detection survives delay only when timers are set above the round-trip	Every tool assumes replies return inside a timeout set for terrestrial latency; once the round-trip exceeds it, answers arrive too late to be used
A timeout below the round-trip does not degrade gracefully; detection is lost or replaced by errors	Below the RTT no reply can return in time, so a listing becomes a timeout, an error, or a failed handshake rather than a partial result
Scan time rises sharply under delay	The scan blocks on network replies, so each message waits at least one round-trip, and dependent lookups wait several
The filter’s own lookups dominate the link, far above the mail payload	Network rules issue many DNS or catalogue queries per message, while the email itself is small by comparison
Delay changes lookup volume, not just speed	A per-attempt timeout below the RTT retransmits unanswered queries, so traffic is spent on lookups that mostly time out

Table 5.18: Behaviour common to SpamAssassin, rspamd, and Razor under delay.

6 | Conclusions and Future Work

6.1 Conclusion

This dissertation tested whether conventional anti-spam tools work over a delay-tolerant network. A testbed was implemented using ION-DTN, BMail gateway, Postfix/Dovecot, local unbound DNS resolver, and `tc netem` to add delay. Three different tools – SpamAssassin, rspamd, and Vipul’s Razor were deployed and evaluated against an earth-to-moon delay scenario with a non-delayed setup as the baseline. The conclusion is that the three filters mentioned above do not tolerate delay by design. They operate under assumptions that round-trips are finished within some timeout values set for normal terrestrial latencies. Once the round trip exceeds that value, the filters become less efficient or fail entirely.

However, spam detection is possible with appropriate configuration of the resolver. By default, the resolver will fail silently without any indication: the filter will still check the messages and return a score but will not detect spam at all. This is the most problematic type of failure because nothing will indicate it. If the tool’s own timeout is set below the round-trip time, detection does not degrade gracefully, it is lost outright: blocklist detections are replaced by timeouts, and a tool that relies on a single collaborative server fails completely.

Delay also makes scanning slow. A scan is blocked until all responses from its network appear, therefore, instead of a sub-second scan, it takes several seconds. The cost depends on the depth of dependent queries, not the number of them: a message

where lookups are dependent on each other gets blocked for the sum of round trips. Filter on the point of delivery is not efficient. The mail queue gets built behind messages that wait for the link. The throughput is limited by the delay, not by the server.

The result is the following - network traffic: DNS lookups of the filter dominate the link way more than the actual email filtering does. On a constrained deep-space link, the anti-spam lookup, not email, becomes a primary consumer of bandwidth. Delay does not just slow filtering down, it changes how much traffic crosses the link, and each tool reacts differently. In rspamd, a DNS timeout shorter than the round trip causes every unanswered query to be resent, so query volume rises for no benefit. In SpamAssassin, the resolver eventually collapses and stops sending queries, so traffic drops instead. In Razor, the load is not DNS at all but repeated TCP connection attempts. In every case, bandwidth is spent on lookups that mostly time out.

These results were measured at Earth–Moon delay. Earth–Mars delay is minutes each way instead of seconds, so the round-trip is far larger than any timeout in the stack. At Mars distance, no default timeout is above the round-trip, so the silent-collapse and total-failure cases seen in the reduced-timeout runs become the normal outcome for all three tools. Dependency-chain scans, which took seconds at Moon distance, would take tens of minutes per message. Retransmission would push DNS traffic to its maximum on a link with little spare capacity. Synchronous, DNS-based anti-spam cannot be used at interplanetary distance without changes.

These tools work over DTN only if every timeout in the stack is set above the link delay and the extra DNS traffic is accepted. Beyond Moon distance this is not enough, and the filtering has to be changed to work with store-and-forward instead of real-time lookups.

6.2 Limitations

This study met its objectives, but is subject to the following limitations:

- **Single fixed delay.** All the measurements are done using only one round-trip delay between Earth and Moon. Larger delays, as well as delay with jitter, packet losses, or broken links, are not tested. The claims on the distance between Earth and Mars are purely logical conclusions based on the mechanisms, but not measured facts.
- **Drift of blocklist and interval between runs.** Since blocklists evolve with time, comparison of the verdict for each message between different runs, not performed at the same time, involves both delay and blocklist drift. Runs 2 and 3 are performed two days apart, so per-message flips test is influenced by this drift. The per-run total numbers of flips can be considered pure measurements.
- **Delay only, not a full space link.** The delay was set up using `tc netem` in a one-node ION-DTN testbed. It simulates the high round-trip time correctly but not any other characteristic of a real deep-space link, like scheduled periods of communication, prolonged interruptions and error-prone transmissions. The findings depict the effects of the delay, not those of the intermittent connection.
- **Behavior is version-dependent.** The outcomes are tied to the exact versions of the programs used. For instance, it is impossible to limit the resolver collapse of unbound because there is no such setting as `infra-cache-max-rtt` in unbound 1.13.1. It was introduced in a newer version. Hence, the findings apply to the specific versions of the tools.

6.3 Future Work

The findings and limitations of this study point to the following future work:

- **Test at Mars-scale and under challenging conditions.** The facts about distance from Earth to Mars in this work are an inference drawn from the mechanisms, and not empirical findings. Testing the tools at minute-scale latency cycles,

along with jitter, packet loss, and disruption in addition to latency, will determine if the tools act according to the mechanism.

- **Move filtering off the synchronous per-message path.** The tested tools do their check online, consulting the blocklist servers along the path for each message, and it causes the performance problems seen here. There are two ways to avoid this approach. In the first one, the blocklist database is moved to the receiver (Moon) side so that per-message check would be done locally and instantly, while a separate process will be doing periodic syncing of the database across the path asynchronously. In the second one, the check is done on the Earth side where blocklist servers are located with normal latency, and the result of the check is added to the email message. This way, only the result of the check crosses the long path. The first approach requires no trusted upstream but allows some staleness of the database; the second one uses up-to-date blocklists but requires a trusted checkpoint on the Earth side.
- **Deferred, asynchronous filtering.** A minor change would be to retain the live lookups, but to avoid them delaying delivery. The message is accepted and saved at receipt time, the DNS or catalog lookups are carried out asynchronously, and the result applied once the lookup completes, potentially minutes afterwards. The message is not delayed for the round-trip delay, and the result is delivered using hold and release or reclassification after delivery. This is in line with the store-and-forward nature of DTN, and much more immediately implementable than redesigning the data path.

Bibliography

- [1] Leigh Torgerson, Scott C. Burleigh, Howard Weiss, Adrian J. Hooke, Kevin Fall, Dr. Vinton G. Cerf, Keith Scott, and Robert C. Durst. Delay-Tolerant Networking Architecture. RFC 4838, April 2007. URL <https://www.rfc-editor.org/info/rfc4838>.
- [2] Scott Burleigh, Kevin Fall, and Edward J. Birrane. Bundle Protocol Version 7. RFC 9171, January 2022. URL <https://www.rfc-editor.org/info/rfc9171>.
- [3] Yosuke Kaneko, Vinton Cerf, Scott Pace, Jim Green, Laura DeNardis, James Schier, Dave J. Israel, Felix Flentge, Leigh Torgerson, Scott Burleigh, Marc Blanchet, Ed Birrane, Keith Scott, Oscar Garcia, Kiyohisa Suzuki, Marius Feldmann, Laura Chappell, Ginny Spicer, Henry Danielson, and Helen Tabunshchyk. Report: Solar System Internet Architecture and Governance. Technical report, Interplanetary Chapter, Internet Society, September 2023.
- [4] Scott Burleigh. Interplanetary overlay network: An implementation of the dtn bundle protocol. In *2007 4th IEEE Consumer Communications and Networking Conference*, pages 222–226, 2007. doi: 10.1109/CCNC.2007.51.
- [5] Jose Velazco. An inter planetary network enabled by smallsats. In *2020 IEEE Aerospace Conference*, pages 1–10, 2020. doi: 10.1109/AERO47225.2020.9172696.
- [6] Matthew Cosby and Wallace Tai. Lunar Communications Architecture Study Report. Technical report, Interagency Operations Advisory Group (IOAG), January 2022. URL [https://www.ioag.org/Public%20Documents/Lunar%](https://www.ioag.org/Public%20Documents/Lunar%20Communications%20Architecture%20Study%20Report.pdf)

20communications%20architecture%20study%20report%20FINAL%20v1.3.pdf.
Accessed 2026-07-07.

- [7] Scott M. Johnson. A method for delivery of SMTP messages over Bundle Protocol networks. Internet-Draft draft-johnson-dtn-interplanetary-smtp-01, Internet Engineering Task Force, March 2025. URL <https://datatracker.ietf.org/doc/draft-johnson-dtn-interplanetary-smtp/01/>. Work in Progress.
- [8] Marc Blanchet. Encapsulation of Email over Delay-Tolerant Networks(DTN) using the Bundle Protocol. Internet-Draft draft-blanchet-dtn-email-over-bp-06, Internet Engineering Task Force, September 2025. URL <https://datatracker.ietf.org/doc/draft-blanchet-dtn-email-over-bp/06/>. Work in Progress.
- [9] Lance Batac, Brian Contreras, Adrian Flores Aquino, Jiahao Li, Sophia Liwag, Nathan Luu, Nicholas Martin, Antonio Ortega Guerrero, Anderson Godo Pena Reyes, Ryan Torrez, and Jilei Zou. BPMail: Utilities for Transferring Email over Bundle Protocol. <https://github.com/250MHz/bpmail>, 2025. Accessed 2026-07-07.
- [10] A. Pentland, R. Fletcher, and A. Hasson. Daknet: rethinking connectivity in developing nations. *Computer*, 37(1):78–83, 2004. doi: 10.1109/MC.2004.1260729.
- [11] S. Guo, M. H. Falaki, E. A. Oliver, S. Ur Rahman, A. Seth, M. A. Zaharia, and S. Keshav. Very low-cost internet access using kiosknets. *SIGCOMM Comput. Commun. Rev.*, 37(5):95–100, October 2007. ISSN 0146-4833. doi: 10.1145/1290168.1290181. URL <https://doi.org/10.1145/1290168.1290181>.
- [12] Shrikant Naidu, Suresh Chintada, Moushumi Sen, and Seshadri Raghavan. Challenges in deploying a delay tolerant network. pages 65–72, 09 2008. doi: 10.1145/1409985.1409998.

- [13] Murray Kucherawy and Elizabeth Zwicky. Domain-based Message Authentication, Reporting, and Conformance (DMARC). RFC 7489, March 2015. URL <https://www.rfc-editor.org/info/rfc7489>.
- [14] Scott M. Johnson. An Interplanetary DNS Model. Internet-Draft draft-johnson-dtn-interplanetary-dns-04, Internet Engineering Task Force, March 2025. URL <https://datatracker.ietf.org/doc/draft-johnson-dtn-interplanetary-dns/04/>. Work in Progress.
- [15] Emmanuel Gbenga Dada, Joseph Stephen Bassi, Haruna Chiroma, Shafi'i Muhammad Abdulhamid, Adebayo Olusola Adetunmbi, and Opeyemi Emmanuel Ajibuwa. Machine learning for email spam filtering: review, approaches and open research problems. *Heliyon*, 5(6):e01802, 2019. ISSN 2405-8440. doi: <https://doi.org/10.1016/j.heliyon.2019.e01802>. URL <https://www.sciencedirect.com/science/article/pii/S2405844018353404>.
- [16] Ernesto Damiani, Sabrina Vimercati, Stefano Paraboschi, and Pierangela Samarati. An open digest-based technique for spam detection. volume 2004, pages 559–564, 01 2004.
- [17] The Spamhaus Project. DNSBL Usage and Blocklist Documentation. <https://www.spamhaus.org/blocklists/>, . Accessed 2026-07-07.
- [18] Apache SpamAssassin Project. Mail::SpamAssassin::Conf — SpamAssassin Configuration File. https://spamassassin.apache.org/full/3.4.x/doc/Mail_SpamAssassin_Conf.html, . Accessed 2026-07-07.
- [19] Apache SpamAssassin Project. Mail::SpamAssassin::Plugin::URIDNSBL — Look Up URLs Against DNS Blocklists. https://spamassassin.apache.org/full/3.4.x/doc/Mail_SpamAssassin_Plugin_URIDNSBL.html, . Accessed 2026-07-07.

- [20] Scott Kitterman. Sender Policy Framework (SPF) for Authorizing Use of Domains in Email, Version 1. RFC 7208, April 2014. URL <https://www.rfc-editor.org/info/rfc7208>.
- [21] Murray Kucherawy, Dave Crocker, and Tony Hansen. DomainKeys Identified Mail (DKIM) Signatures. RFC 6376, September 2011. URL <https://www.rfc-editor.org/info/rfc6376>.
- [22] Vsevolod Stakhov. Rspamd: Fast, Free and Open-Source Spam Filtering System. <https://rspamd.com/>. Accessed 2026-07-24.
- [23] Vipul's Razor / Cloudmark. Vipul's Razor: Distributed, Collaborative Spam Detection and Filtering Network. <https://razor.sourceforge.net/>. Accessed 2026-07-07.
- [24] Stephen Farrell, Susan Symington, Howard Weiss, and Peter Lovell. Delay-Tolerant Networking Security Overview. Internet-Draft draft-irtf-dtnrg-sec-overview-06, Internet Engineering Task Force, March 2009. URL <https://datatracker.ietf.org/doc/draft-irtf-dtnrg-sec-overview/06/>. Work in Progress.
- [25] Sushant Jain, Kevin Fall, and Rabin Patra. Routing in a delay tolerant network. *SIGCOMM Comput. Commun. Rev.*, 34(4):145–158, August 2004. ISSN 0146-4833. doi: 10.1145/1030194.1015484. URL <https://doi.org/10.1145/1030194.1015484>.
- [26] Marc Blanchet. Encapsulation of HTTP over Delay-Tolerant Networks(DTN) using the Bundle Protocol. Internet-Draft draft-blanchet-dtn-http-over-bp-04, Internet Engineering Task Force, September 2025. URL <https://datatracker.ietf.org/doc/draft-blanchet-dtn-http-over-bp/04/>. Work in Progress.

- [27] The Spamhaus Project. QNAME Minimization and Spamhaus DNSBLs — an Explanation. <https://www.spamhaus.org/resource-hub/dnsbl/qname-minimization-and-spamhaus-dnsbls/>, . Accessed 2026-07-24.
- [28] Paul Shupak. Network Tests Latency. <https://cwiki.apache.org/confluence/spaces/SPAMASSASSIN/pages/119544790/NetworkTestsLatency>, 2006. Message to the SpamAssassin users mailing list, 14 January 2006; archived on the Apache SpamAssassin wiki. Accessed 2026-07-24.
- [29] Alireza Pour. Minimizing the time of spam mail detection by relocating filtering system to the sender mail server. *International Journal of Network Security Its Applications*, 4:53–62, 03 2012. doi: 10.5121/ijnsa.2012.4204.
- [30] NASA Jet Propulsion Laboratory. Interplanetary Overlay Network (ION-DTN). <https://sourceforge.net/projects/ion-dtn/>. Version 4.1.4. Accessed 2026-07-07.
- [31] Arch Linux. Unbound — archwiki. URL <https://wiki.archlinux.org/title/Unbound>. Accessed 2026-07-02.
- [32] NLnet Labs. Unbound timeout behaviour, 2025. URL <https://www.nlnetlabs.nl/documentation/unbound/info-timeout/>.

A1 | Appendix

A1.1 Use of Artificial Intelligence

This appendix declares the use of generative AI tools in this dissertation, in line with the College's policy on academic integrity.

The tool used was Claude (Anthropic). Model: Claude Opus 4.8. Accessed through the web interface over the course of the project. It was used as an assistant in debugging.

When a build or component was failing in a non-obvious way, this tool would be employed to analyze the failure and to identify its source. The most obvious use-case for such analysis was rspamd crashing on the ARM64 test machine - rspamd terminated because of illegal instruction trap, and this tool helped to identify that the problem was caused by the host incorrectly reporting some CPU feature (SVE2), and consequently, by the shim required to run rspamd on this architecture. Another use of this kind was analysis of unbound resolver's behavior with delay and analysis of failing builds of DTN and mail layers - in all cases this tool gave a hint on where to look for the solution, and the solution was validated by testing it on a live system.